



K.R. MANGALAM UNIVERSITY
THE COMPLETE WORLD OF EDUCATION

OPERATING SYSTEM CAPSTONE PROJECT

Submitted in partial fulfilment of the requirement of the degree

BACHELORS OF TECHNOLOGY

In

CSE With Specialization (AI & ML)

To

K.R Mangalam University

By

Vaibhav Kedia (2301730096)

Under the supervision of

Mr. Mohammad Aijaz



School of Engineering and Technology
K.R Mangalam University, Gurugram- 122103, India

Part A: Fundamental Concepts

Question 1: Why Operating Systems Matter Despite Hardware Advances

Simplified Answer:

Even though modern hardware is very fast, it's still difficult for software applications to work directly with it. An operating system (OS) acts as a bridge between hardware and software. It makes things simpler and safer.

Key Services of an OS:

OS Service	What It Does	Example
Process Management	Creates, schedules, and manages programs while they run	Lets 10 applications run at the same time by sharing CPU time
Memory Management	Allocates and manages RAM; keeps programs from crashing into each other's memory	Each program gets its own memory space so they don't interfere
I/O Management	Controls communication between programs and devices (disk, printers, network)	When you save a file, the OS controls writing data to the hard drive

Real-World Example:

When you run 10 applications at once, the OS does three things:

- Divides the CPU time among them using scheduling algorithms
- Gives each program its own virtual memory space
- Manages input/output requests to devices

Question 2: Comparing OS Structures - Monolithic vs Layered vs Microkernel

Simplified Comparison:

Feature	Monolithic	Layered	Microkernel
Structure	All services in kernel (big single program)	Services divided into organized layers	Only essential services in kernel; others in user space
Advantages	Fast performance	Better organization and maintainability	High reliability and security
Disadvantages	If one service crashes, entire OS crashes; hard to fix	More layers = slower performance	More complex; requires more communication

Best Choice for Distributed Web Applications: Microkernel

Why Microkernel?

1. **Reliability:** If a device driver or file service crashes, the core OS doesn't fail—only that service stops
2. **Security:** Less code runs in kernel mode (dangerous), so fewer security risks
3. **Fault Tolerance:** Web applications need to keep running even if parts fail
4. **Scalability:** Services can be updated without restarting the entire system

Example: If a network driver crashes in a Microkernel, your web server keeps running. In a Monolithic OS, the entire system would crash.

Question 3: Are Threads More Efficient Than Processes?

Critical Analysis:

This statement is **generally TRUE**, but with important exceptions.

Why Threads Are More Efficient:

Aspect	Processes	Threads
Memory Usage	Each has separate memory space (heavy)	Share memory with their process (light)
Context Switching	Slow (save/restore entire memory and state)	Fast (save/restore only CPU registers)
Creation Time	Slow	Fast
Communication	Requires Inter-Process Communication (IPC) - complex	Direct shared memory - simple
CPU Utilization	Multi-processing increases CPU use	Multi-threading increases CPU use more efficiently

Why Threads Can Be Problematic:

1. **Synchronization Issues:** Threads share memory, so they need locks to prevent data corruption
2. **Crash Risk:** If one thread crashes badly, it can crash the entire process (affecting all threads)
3. **Debugging Difficulty:** Hard to track bugs in multi-threaded programs

Conclusion: Threads are more efficient for most tasks, but processes provide better isolation and safety.

Question 4: Memory Allocation - First-Fit vs Best-Fit

Problem Setup:

- Processes need: 12 MB, 18 MB, 6 MB
- Available memory blocks: 20 MB, 10 MB, 15 MB

First-Fit Allocation:

Allocate each process to the first block large enough.

Process	Required	Block Chosen	Leftover
P1	12 MB	20 MB	8 MB remaining
P2	18 MB	15 MB	Not enough! (block too small)
P3	6 MB	10 MB	4 MB remaining

Result: P2 cannot be allocated. Free fragments: 8 MB, 10 MB, 4 MB (Total wasted: 22 MB)

Best-Fit Allocation:

Allocate each process to the smallest block that fits (to minimize leftover space).

Process	Required	Best Block	Leftover
P1	12 MB	15 MB	3 MB remaining
P2	18 MB	20 MB	2 MB remaining
P3	6 MB	10 MB	4 MB remaining

Result: All processes allocated! Free fragments: 3 MB, 2 MB, 4 MB (Total wasted: 9 MB)

Fragmentation Analysis:

- **First-Fit:** Creates larger fragments (8 MB, 10 MB, 4 MB) - harder to reuse
- **Best-Fit:** Creates smaller fragments (3 MB, 2 MB, 4 MB) - better fragmentation but slower search

Conclusion: Best-Fit performs better here, but First-Fit is usually faster in practice

Part B: Scheduling & Synchronization

Question 5: CPU Scheduling Algorithms Comparison

Given Process Data:

- P1: BT = 5, AT = 0
- P2: BT = 3, AT = 1
- P3: BT = 8, AT = 2
- P4: BT = 6, AT = 3

CODE FOR FCFS,SJF(non- preemptive),RR(QT=4ms):

```

from collections import deque
# PROCESSES: (name, burst_time, arrival_time)
PROCESSES = [
    ("P1", 5, 0),
    ("P2", 3, 1),
    ("P3", 8, 2),
    ("P4", 6, 3)
]
def fcfs(processes):
    t = 0
    result = []
    for name, bt, at in sorted(processes, key=lambda x: x[2]):
        if t < at:
            t = at
        start = t
        finish = t + bt
        t = finish
        result.append((name, start, finish, at, bt))
    return result
def sjf_non_preemptive(processes):
    t = 0
    remaining = processes[:]
    result = []
    while remaining:
        available = [p for p in remaining if p[2] <= t]
        if not available:
            t = min(p[2] for p in remaining)
        available = [p for p in remaining if p[2] <= t]
        name, bt, at = min(available, key=lambda x: x[1])
        start = t
        finish = t + bt
        t = finish
        result.append((name, start, finish, at, bt))
        remaining.remove((name, bt, at))
    return result
def round_robin(processes, quantum=4):
    procs = sorted(processes, key=lambda x: x[2])
    n = len(procs)
    remaining_bt = [bt for _, bt, _ in procs]
    t = 0
    i = 0
    q = deque()
    completed = [False] * n
    completion_time = [0] * n
    gantt = []
    while True:
        while i < n and procs[i][2] <= t:
            q.append(i)
            i += 1
        if not q:
            if i < n:
                t = procs[i][2]
                continue
            else:
                break
        idx = q.popleft()
        name, bt, at = procs[idx]
        run = min(quantum, remaining_bt[idx])
        start = t
        t += run
        finish = t
        gantt.append((name, start, finish))
        remaining_bt[idx] -= run
        while i < n and procs[i][2] <= t:
            q.append(i)
            i += 1
        if remaining_bt[idx] > 0:
            q.append(idx)
        else:
            completed[idx] = True
            completion_time[idx] = t
        if all(completed):
            break
    result = [(procs[idx][0], procs[idx][2], completion_time[idx], procs[idx][1]) for idx in range(n)]
    return gantt, result
def compute_metrics(result):
    wt, tat = [], []
    for name, at, ct, bt in result:
        turnaround = ct - at
        waiting = turnaround - bt
        wt.append(waiting)
        tat.append(turnaround)
    return wt, tat, sum(wt)/len(wt), sum(tat)/len(tat)
# RUN ALL ALGORITHMS
print("== FCFS ==")
fcfs_res = fcfs(PROCESSES)
print("Gantt: " + "\n".join([f"[{name}:{start}-{finish}]" for name, start, finish, _, _ in fcfs_res]))
wt, tat, avg_wt, avg_tat = compute_metrics([(r[0], r[3], r[2], r[1]) for r in fcfs_res])
print(f"Avg WT: {avg_wt:.2f}, Avg TAT: {avg_tat:.2f}\n")
print("== SJF ==")
sjf_res = sjf_non_preemptive(PROCESSES)
print("Gantt: " + "\n".join([f"[{name}:{start}-{finish}]" for name, start, finish, _, _ in sjf_res]))
wt, tat, avg_wt, avg_tat = compute_metrics([(r[0], r[3], r[2], r[1]) for r in sjf_res])
print(f"Avg WT: {avg_wt:.2f}, Avg TAT: {avg_tat:.2f}\n")
print("== ROUND ROBIN (q=4) ==")
rr_gantt, rr_res = round_robin(PROCESSES)
print("Gantt: " + "\n".join([f"[{name}:{start}-{finish}]" for name, start, finish in rr_gantt]))
wt, tat, avg_wt, avg_tat = compute_metrics(rr_res)
print(f"Avg WT: {avg_wt:.2f}, Avg TAT: {avg_tat:.2f}")

```

a) Gantt Charts

FCFS (First-Come-First-Served):

Timeline: [P1: 0-5] [P2: 5-8] [P3: 8-16] [P4: 16-22]

SJF (Shortest Job First - Non-preemptive):

Timeline: [P1: 0-5] [P2: 5-8] [P4: 8-14] [P3: 14-22]

Round Robin (Quantum = 4ms):

Timeline: [P1: 0-4] [P2: 4-7] [P3: 7-11] [P4: 11-15] [P3: 15-19] [P4: 19-22]

b) Average waiting and turnaround times

Output:

```
==== FCFS ====
Gantt: [P1:0-5] [P2:5-8] [P3:8-16] [P4:16-22]
Avg WT: 4.00, Avg TAT: 11.25

==== SJF ====
Gantt: [P1:0-5] [P2:5-8] [P4:8-14] [P3:14-22]
Avg WT: 4.00, Avg TAT: 10.75

==== ROUND ROBIN (q=4) ====
Gantt: [P1:0-4] [P2:4-7] [P3:7-11] [P4:11-15] [P1:15-16] [P3:16-20] [P4:20-22]
Avg WT: 9.25, Avg TAT: 14.75
```

c) Algorithm Comparison

Best for Multiprogrammed Systems: SJF

Algorithm	Throughput	Fairness	Waiting Time
FCFS	Lower	Fair	High for long jobs
SJF	Highest	Unfair (long jobs wait)	Lowest
Round Robin	Medium	Fair	Medium

Recommendation: Use **SJF** for batch systems (best throughput) or **Round Robin** for interactive systems (better fairness).

Question 6: Deadlock Handling in Banking Systems

a) Banker's Algorithm - Deadlock Avoidance

How It Works (Simple Explanation):

1. Each transaction declares maximum resources it needs upfront
2. When requesting resources, the OS asks: "If I give this resource, will the system still be safe?"
3. "Safe" means all transactions can eventually complete without deadlock
4. Only grant requests that keep the system in a safe state

Example: Bank transfers involving multiple accounts

- Transaction T1 needs accounts (A, B, C) - declares max need upfront
- T1 requests to lock account A - OS checks: "Can T1 and others finish if A is locked?"
- If YES → Lock granted
- If NO → Lock delayed

b) Deadlock Detection & Recovery in Programming

Detection Mechanism:

The OS periodically checks for cycles in the "Wait-for Graph":

- If P1 waits for P2, P2 waits for P3, P3 waits for P1 → **Cycle = Deadlock!**

Recovery Approach:

```
from collections import defaultdict, deque

def detect_deadlock(wait_for):
    """Detect cycle in wait-for graph using Kahn's algorithm"""
    graph = defaultdict(list)
    indegree = defaultdict(int)

    # Build graph and indegree
    for p, waits in wait_for.items():
        for q in waits:
            graph[p].append(q)
            indegree[q] += 1
            if p not in indegree:
                indegree[p] = 0

    # BFS with nodes having indegree 0
    q = deque([p for p in indegree if indegree[p] == 0])
    visited = 0

    while q:
        node = q.popleft()
        visited += 1
        for neighbor in graph[node]:
            indegree[neighbor] -= 1
            if indegree[neighbor] == 0:
                q.append(neighbor)

    deadlock = visited < len(indegree)
    print("DEADLOCK DETECTED!" if deadlock else "No deadlock")
    return deadlock

def recover(wait_for, victim):
    """Abort victim transaction"""
    if victim in wait_for:
        del wait_for[victim]
    for p in list(wait_for.keys()):
        if victim in wait_for[p]:
            wait_for[p].remove(victim)
    print(f"RECOVERED: Aborted {victim}")

# TEST CASE: Deadlock (T1→T2→T3→T1)
wait_for = {"T1": ["T2"], "T2": ["T3"], "T3": ["T1"]}
print("Initial state:")
detect_deadlock(wait_for)
recover(wait_for, "T3")
print("\nAfter recovery:")
detect_deadlock(wait_for)
```

Output:

Initial state:

DEADLOCK DETECTED!

RECOVERED: Aborted T3

After recovery:

No deadlock

Recovery Strategies:

1. **Abort Victim:** Stop one transaction, release its resources
2. **Rollback:** Undo partial work and restart
3. **Preempt Resources:** Force one transaction to release locks
4. **Restart Later:** Put transaction back in queue for retry

Banking Example: If T1 and T2 deadlock over accounts:

- Abort T2 (cheaper operation)
- Release all locks held by T2
- Rollback T2's changes
- Restart T2 when system is free

Question 7: Producer-Consumer with Semaphores

Problem: Producer and Consumer share a buffer. Need to prevent race conditions.

Simple Explanation:

- **Producer** creates data and puts it in buffer
- **Consumer** takes data from buffer and uses it
- **Problem:** Both might access buffer simultaneously, corrupting data
- **Solution:** Use semaphores (special counters) to synchronize

CODE:

```
import threading
import time
import random

BUFFER_SIZE = 5
buffer = []
empty = threading.Semaphore(BUFFER_SIZE)
full = threading.Semaphore(0)
mutex = threading.Semaphore(1)

def producer():
    for i in range(10):
        time.sleep(random.uniform(0.1, 0.3))
        item = f"Item-{i}"

        empty.acquire()
        mutex.acquire()
        buffer.append(item)
        print(f"Producer added {item} | Buffer: {len(buffer)}/{BUFFER_SIZE}")
        mutex.release()
        full.release()

def consumer():
    for i in range(10):
        time.sleep(random.uniform(0.2, 0.5))

        full.acquire()
        mutex.acquire()
        item = buffer.pop(0)
        print(f"Consumer removed {item} | Buffer: {len(buffer)}/{BUFFER_SIZE}")
        mutex.release()
        empty.release()

# RUN
print("Starting Producer-Consumer...")
p_thread = threading.Thread(target=producer)
c_thread = threading.Thread(target=consumer)

p_thread.start()
c_thread.start()

p_thread.join()
c_thread.join()
print("COMPLETED!")
```

Output:

```
Starting Producer-Consumer...
Producer added Item-0 | Buffer: 1/5
Producer added Item-1 | Buffer: 2/5
Consumer removed Item-0 | Buffer: 1/5
Producer added Item-2 | Buffer: 2/5
Producer added Item-3 | Buffer: 3/5
Consumer removed Item-1 | Buffer: 2/5
Producer added Item-4 | Buffer: 3/5
Producer added Item-5 | Buffer: 4/5
Consumer removed Item-2 | Buffer: 3/5
Producer added Item-6 | Buffer: 4/5
Consumer removed Item-3 | Buffer: 3/5
Producer added Item-7 | Buffer: 4/5
Producer added Item-8 | Buffer: 5/5
Consumer removed Item-4 | Buffer: 4/5
Producer added Item-9 | Buffer: 5/5
Consumer removed Item-5 | Buffer: 4/5
Consumer removed Item-6 | Buffer: 3/5
Consumer removed Item-7 | Buffer: 2/5
Consumer removed Item-8 | Buffer: 1/5
Consumer removed Item-9 | Buffer: 0/5
COMPLETED!
```

How Semaphores Prevent Race Conditions:

Semaphore	Purpose	Value Change
empty	Counts empty slots in buffer	Decreases when producer adds; increases when consumer takes
full	Counts filled slots in buffer	Increases when producer adds; decreases when consumer takes
mutex	Ensures only one thread accesses buffer at a time	Acquired when entering critical section; released after

Safety Guarantee: Only one thread accesses the buffer at a time, so no corruption!

Question 8: Page Replacement - FIFO vs LRU

Given Page Access Sequence: [2, 1, 4, 2, 3, 4, 3]

Frame Size: 3 (can hold 3 pages at once)

FIFO (First-In-First-Out)

Replace the oldest page in memory.

Step	Page	Frames After	Page Fault?
1	2	[2]	Yes
2	1	[2, 1]	Yes
3	4	[2, 1, 4]	Yes
4	2	[2, 1, 4]	No (already there)
5	3	[1, 4, 3]	Yes (remove 2, oldest)
6	4	[1, 4, 3]	No (already there)
7	3	[1, 4, 3]	No (already there)

Total FIFO Page Faults: 4

LRU (Least Recently Used)

Replace the page not used for the longest time.

Step	Page	Frames After	Page Fault?
1	2	[2]	Yes
2	1	[2, 1]	Yes
3	4	[2, 1, 4]	Yes
4	2	[1, 4, 2]	No (mark 2 as recent)

5	3	[4, 2, 3]	Yes (remove 1, least recent)
6	4	[4, 2, 3]	No (already there)
7	3	[4, 2, 3]	No (already there)

Total LRU Page Faults: 4

Comparison

Algorithm	Page Faults	Best When
FIFO	4	Rarely optimal; simple but not smart
LRU	4	Access patterns have locality (pages used recently will be used again)

Conclusion: LRU generally performs better than FIFO in real systems because programs tend to reuse recently-accessed pages (temporal locality principle).

Part C: Distributed Systems

Question 9: Designing a Distributed File System

a) Two Critical Issues in Distributed OS Design

Issue 1: Consistency in File Sharing

Problem: Files are stored across multiple servers. When two users access the same file from different locations:

- User A modifies the file on Server 1
- User B reads the file from Server 2
- Should User B see User A's changes immediately?

Challenges:

- Network delays: Changes take time to propagate
- Caching: Each server caches files locally for speed
- Conflicts: What if both users modify the same file simultaneously?

Issue 2: Resource Management & Load Balancing

Problem: Multiple users and devices share limited server resources.

Challenges:

- One server might become overloaded (too many requests)
- Other servers might be underutilized

- Need automatic load balancing to distribute work
- Failure handling: If one server crashes, work must move elsewhere

b) Architectural Approaches for Distributed File Management

Approach	How It Works	Benefits
Client-Server with Replication	Files replicated across multiple servers; clients read from nearest/fastest server	Improves performance and availability; if one server fails, copies exist elsewhere
Distributed Lock Manager (DLM)	Central lock service coordinates file access across all servers; prevents simultaneous conflicting modifications	Prevents data corruption when multiple clients access same file
Caching with Coherence Protocol	Servers cache files locally but maintain consistency using coherence protocol (similar to CPU caches)	Fast access through caching; consistency through protocol

Recommended Architecture: Combine all three:

1. **Replicate** files across multiple data centers
2. **Use DLM** to manage concurrent access
3. **Cache** aggressively with coherence protocol

This ensures **performance, consistency, and fault tolerance**.

Question 10: Synchronous vs Asynchronous Checkpointing

What is Checkpointing? Periodically save system state so recovery is possible after failure.

Synchronous Checkpointing

How It Works:

1. Coordinator tells all nodes: "Stop! Checkpoint now!"
2. All nodes freeze current execution
3. All nodes save their state to stable storage
4. All nodes synchronize timestamps
5. All nodes resume execution

Visual Example:

Time: 0s 1s 2s 3s 4s 5s

Node 1: [Running] [STOP] [Save] [Sync] [Resume] [Running]

Node 2: [Running] [STOP] [Save] [Sync] [Resume] [Running]

Node 3: [Running] [STOP] [Save] [Sync] [Resume] [Running]

←→ All synchronized →←

Strengths:

- **Global Consistency:** All checkpoints aligned - safe recovery to any checkpoint
- **Easy Recovery:** All nodes at same state; simple rollback
- **Transparent:** Application doesn't need to manage checkpoints

Weaknesses:

- **High Blocking Time:** All nodes stop (entire system pauses)
- **Scalability Issue:** With many nodes, coordination overhead grows
- **Performance Penalty:** Frequent synchronization slows everything

Asynchronous Checkpointing

How It Works:

1. Each node saves state independently, whenever it wants
2. No global synchronization
3. Nodes continue running without waiting for others

Visual Example:

Time: 0s 1s 2s 3s 4s 5s

Node 1: [Running] [Save] [Running] [Save] [Running]

Node 2: [Running] [Running] [Save] [Running] [Save]

Node 3: [Save] [Running] [Running] [Save] [Running]

←— Each independent —→

Strengths:

- **No Blocking:** Nodes never stop; continuous progress
- **High Performance:** Minimal overhead
- **Scalability:** Works well with many nodes
- **Better for IoT/Edge:** Devices can checkpoint independently

Weaknesses:

- **Recovery Complexity:** Checkpoints not aligned; need to reconstruct consistent state
- **Cascading Rollbacks:** Recovering to one node's checkpoint might require rolling back others
- **Data Loss Risk:** If recovery isn't done carefully, some committed work might be lost

Comparison Table

Aspect	Synchronous	Asynchronous
Blocking Time	High	None
Consistency	Easy to ensure	Complex
Recovery	Simple	Complex (may need cascade)
Performance	Slower	Faster
Scalability	Poor	Good
Best For	Small systems, critical applications	Large distributed systems, edge computing

Recommendation:

- Use **Synchronous** for small distributed databases where consistency is critical
- Use **Asynchronous** for large-scale systems like cloud data centers

Question 11: Smart Home IoT Scheduling

Scenario: Smart home with cameras, lights, locks, temperature sensors. Need to prioritize security over comfort.

a) Process Scheduling Strategy

Priority Levels:

1. **Critical (Highest):** Security devices - Camera motion detection, door lock alerts
2. **High:** Safety - Fire/smoke detectors, temperature warnings
3. **Medium:** Regular - Light automation, routine tasks
4. **Low:** Non-critical - Statistics collection, logging

Recommended Algorithm: Priority Preemptive Scheduling

How It Works:

1. Assign priority number to each device's task (1=Critical, 4=Low)
2. Always run the highest priority ready task
3. If high-priority task arrives, pause lower-priority task
4. Return to lower-priority when high-priority completes

Justification:

- **Security first:** If camera detects break-in, it must execute immediately
- **Preemption:** Can interrupt light dimming to handle security event
- **Fairness:** Low-priority tasks still run when no urgent events occur

b) Inter-Process Communication (IPC) Methods for IoT

Best IPC Methods:

IPC Method	Use Case	Why Suitable
MQTT (Publish-Subscribe)	All devices send events to central hub	Lightweight, efficient; perfect for IoT; handles device failures gracefully
Message Queues	Pass security alerts between services	Reliable delivery with persistence; if consumer crashes, message waits
Shared Memory	Fast local communication between nearby modules	Ultra-fast for local sensors on same gateway; minimal latency
REST API	Remote device queries (e.g., check camera feed)	Stateless, scales well; easy for cloud integration

Question 12: Linux System Calls Demo

Objective: Demonstrate OS system calls for file operations and threading.

Linux is selected as the operating system for this case study. The goal is to show how a simple Python program uses Linux system calls indirectly to create, read, inspect, and delete a file.

```
import os

FILENAME = "linux_demo.txt"

DATA = "Hello from a Linux file system call demo in Python!\n"

def main():

    print("== Python File Demo on Linux ==\n")

    # STEP 1: Create and write a file

    # This uses high-level Python I/O, which internally calls Linux sys_open and sys_write.

    print("STEP 1: Creating and writing file...")

    with open(FILENAME, "w", encoding="utf-8") as f:

        f.write(DATA)

    print(f" Wrote text to {FILENAME}\n")

    # STEP 2: Read the file back

    # Internally mapped to Linux sys_open and sys_read.

    print("STEP 2: Reading file...")
```

```

with open(FILENAME, "r", encoding="utf-8") as f:
    content = f.read()
    print(" File content:")
    print(" ", content)

    # STEP 3: Get file metadata using os.stat
    # This wraps the Linux stat() system call.

    print("STEP 3: File metadata (stat):")

    st = os.stat(FILENAME)
    print(f" Size: {st.st_size} bytes")
    print(f" Permissions (octal): {oct(st.st_mode & 0o777)}")
    print(f" Inode number: {st.st_ino}")

    print()

    # STEP 4: Delete the file using os.remove (unlink system call)

    print("STEP 4: Deleting file...")
    os.remove(FILENAME)

    print(" File deleted. Demo complete.")

if __name__ == "__main__":
    main()

```

- Linux provides system calls such as open, read, write, stat, and unlink to allow user programs to work with files. Python's open(), os.stat(), and os.remove() are high-level wrappers that eventually invoke these underlying system calls through the C library, which triggers a switch from user mode to kernel mode.
- This program demonstrates the file abstraction offered by the operating system. The Python code never manipulates disk blocks directly; instead, it treats the file as a simple stream of bytes, while the Linux kernel handles details like locating the inode, managing the file descriptor, caching data, and scheduling disk I/O.
- Security and protection are also illustrated. When open() or os.stat() are called, the kernel checks file permissions (read/write bits in the mode field) before allowing access, enforcing the access control policy defined by the Linux permission model